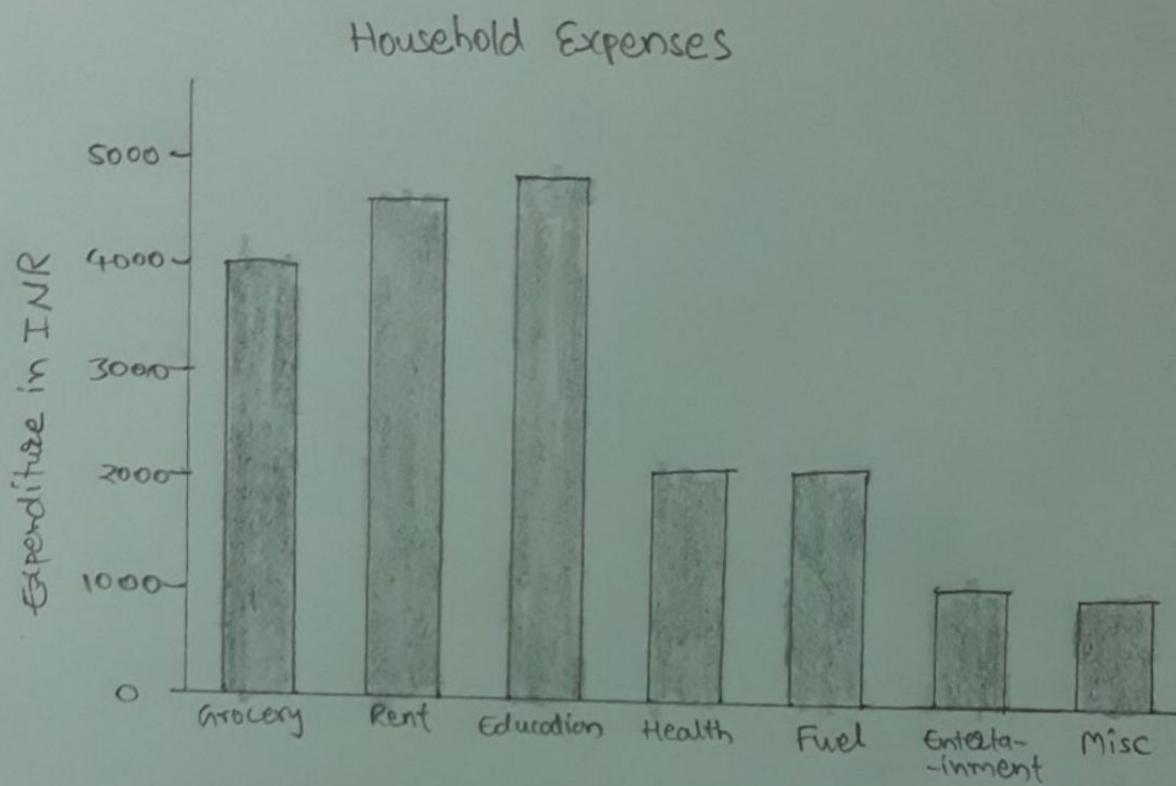


Output -



Machine Learning Lab Using Python

TITLE: Bar Graph.

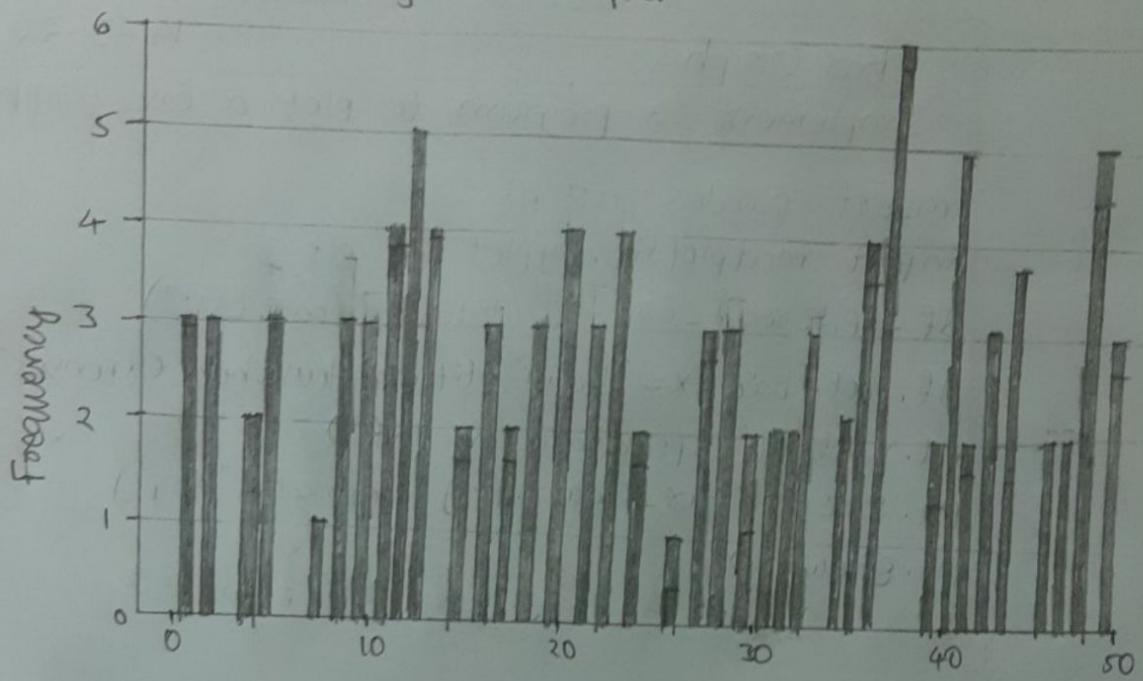
DATE: 14-3-22

Aim: Implement a program to plot a bar graph

```
import pandas as pd  
import matplotlib.pyplot as plt  
df = pd.read_csv("../datasets/bar.csv")  
df.plot.bar(x="Head", title="Household Expenses")  
plt.ylabel("Expenditure in INR")  
plt.savefig("bar-Chart.png", format="png")  
plt.show()
```

Output -

Histogram Example.



TITLE: Histogram

DATE: 14-3-22

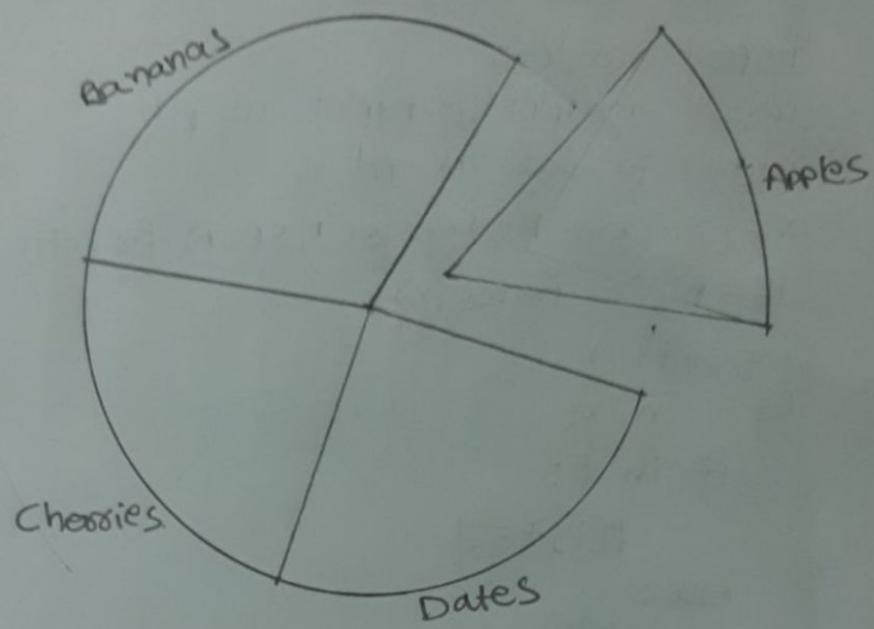
Aim: Implement a program to plot a histogram

```

import random
import matplotlib.pyplot as plt
import pandas as pd
x=[random.randrange(1,50,1) for i in range(100)]
df=pd.DataFrame(x)
f=dict()
for i in x:
    if i in f:
        f[i]+=1
    else:
        f[i]=1
plt.bar(f.keys(),f.values(),color="green")
plt.ylabel("Frequency")
plt.title("Histogram Example")
plt.savefig("Histogram.png",format="png")
plt.show()

```

Output -



DATE: 14-3-22

TITLE: Pie Chart

Aim: Implement a Program to plot a PieChart

```
import matplotlib.pyplot as plt  
import numpy as np  
lst = np.random.sample(4)  
y = np.array(lst)  
mylabels = ["Apples", "Bananas", "Cherries", "Dates"]  
myexplode = [0.2, 0, 0, 0]  
plt.pie(y, labels=mylabels, explode=myexplode,  
        startangle=0)  
plt.savefig("Pie.png", format="png")  
plt.show()
```

Output -

/home/varun/AI-ML/lab/manual

Index(['LOCATION', 'Country', 'INDICATOR', 'Indicator', 'MEASURE',
'Measure', 'INEQUALITY', 'Inequality', 'Unit Code', 'Unit',
'PowerCode Code', 'PowerCode', 'Reference Period Code',
'Reference Period', 'Value', 'Flag Codes', 'Flags'],
dtype='object').

(3292, 17)

	LOCATION	Country	INDICATOR	Flags
0	AUS	Australia	HO-BASE	- - - - Estimated Value
1	AUT	Austria	HO-BASE	- - - - NaN
2	BEL	Belgium	HO-BASE	- - - - NaN
3	CAN	Canada	HO-BASE	- - - - NaN
4	CZE	Czech Republic	HO-BASE	- - - - NaN

(5 rows x 17 columns)

	LOCATION	Country	INDICATOR	Flags
3287	EST	Estonia	WL-TNow	- - - - NaN
3288	ISR	Israel	WL-TNow	- - - - Estimated Value
3289	RUS	Russia	WL-TNow	- - - - Estimated Value
3290	SVN	Slovenia	WL-TNow	- - - - NaN
3291	OECD	OECD-Total	WL-TNow	- - - - NaN

TITLE: Data Cleaning-MetaData

DATE: 21-3-22

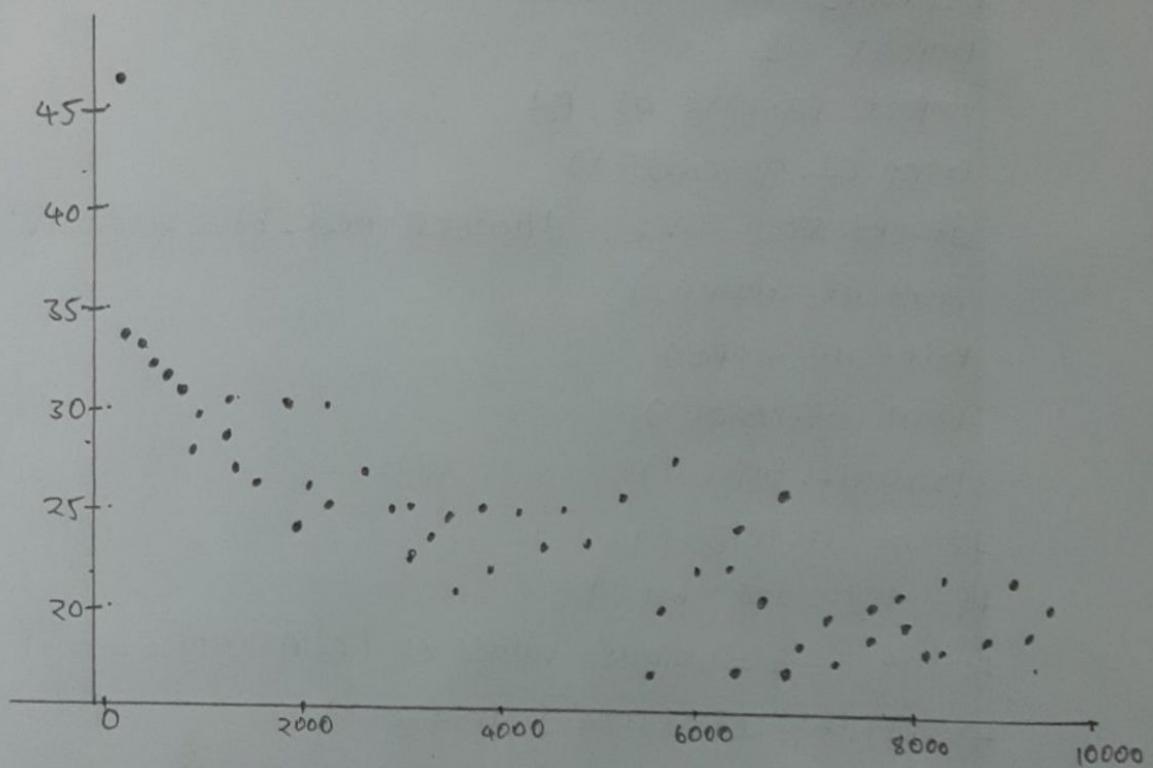
Aim: Implement a program for reading and cleaning datasets.

```

import os
import pandas as pd
print(os.getcwd())
df = pd.read_csv("../datasets/oecd_bli_2015.csv")
print(df.columns)
print(df.shape)
print(df.head())
print(df.tail())
#print(df.info())
#print(df.describe())
#print("== unique values of LOCATION ==")
#print(df[["LOCATION"]].unique())
#print(df[["LOCATION"]].value_counts())
#print(df[["Unit"]].dropna())
#print(df.sort_values("Value"))
#print(df.groupby("LOCATION"))
ndf = df[(df[["Inequality"]] == "Total") &
          (df[["INDICATOR"]] == "HO-BASE")]
#print(ndf)

```

Output -



Number of prime numbers in step of 200 from 0 to 10,000

TITLE: Data Cleaning - Prime Numbers

DATE: 28-3-22.

Aim: Write a program to find number of prime numbers using file handling techniques.

```
import os
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
END = 10000
```

```
LEN = 50
```

```
STEP = int(END/LEN)
```

```
def generate_numbers():
```

```
    with open("numbers.txt", "w") as lfile:
```

```
        lfile.write(str(2) + "\n")
```

```
        for i in range(3, END, 2):
```

```
            lfile.write(str(i) + "\n")
```

```
def is_prime(lnumber):
```

```
    for i in range(2, lnumber):
```

```
        if lnumber % i == 0:
```

```
            return False
```

```
    return True
```

```
def make_data():
```

```
    if os.path.exists("numbers.txt"):
```

```
        os.remove("numbers.txt")
```

```
    generate_numbers()
```

```
    if os.path.exists("primes.txt"):
```

```
        os.remove("primes.txt")
```

```

with open("numbers.txt", "r") as ffile:
    for line in ffile:
        lnumber = int(line.strip("\n"))
        if is_prime(lnumber):
            with open("primes.txt", "a") as pfile:
                primes.write(line)

if __name__ == "__main__":
    make_data()
    prime_range = [STEP * i for i in range(1, (LEN + 1))]
    prime_count = [0 for i in range(1, (LEN + 1))]
    COUNT = 0
    primes = 0
    with open("primes.txt", "r") as file:
        LIMIT = STEP
        for line in file:
            primes += 1
            number = int(line.strip("\n"))
            if number <= LIMIT:
                prime_count[COUNT] += 1
            else:
                LIMIT += STEP
                COUNT += 1
                prime_count[COUNT] += 1

    print(primes)
    plt.scatter(prime_range, prime_count)
    plt.savefig("wlb-plot.png", format="png")
    plt.xticks(prime_range, rotation=45)

```

```
plt.xlabel("numbers")
plt.ylabel("# prime numbers")
plt.show()
```

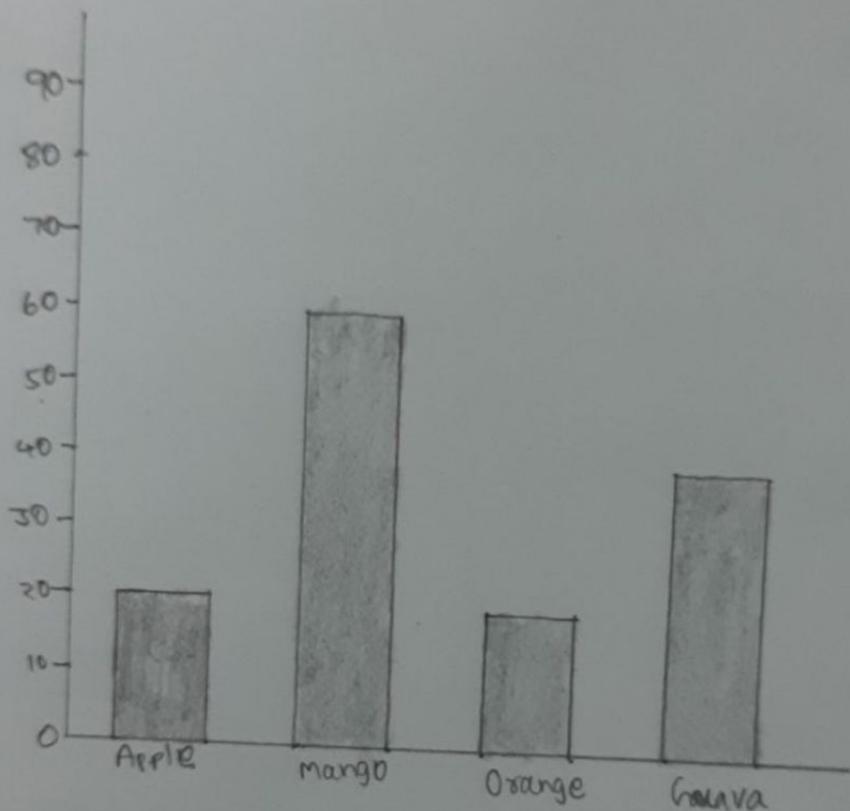
Output -

Frequency of Apple : 20

Frequency of Mango : 60

Frequency of Orange : 20

Frequency of Guava : 40



TITLE: Word Frequency

DATE: 4-4-22

Aim: Write a program to find the word frequency of each individual word in the given text.

```

import Pandas as pd
import matplotlib.pyplot as plt
def myfunc():
    words_dict = {}
    with open("/content/words.txt", "r") as f:
        for line in f:
            words = line.strip("\n").split(" ")
            for word in words:
                if word in words_dict:
                    words_dict[word] += 1
                else:
                    words_dict[word] = 1
    return words_dict
if __name__ == "__main__":
    wd = myfunc()
    for allKeys in wd:
        print("Frequency of ", allKeys, end=" ")
        print(":", end=" ")
        print(wd[allKeys], end=" ")
    print()
    plt.bar(list(wd.keys()), list(wd.values()))
    plt.show()

```

Output -

{'age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6'}
Coefficient : [938.23786125]

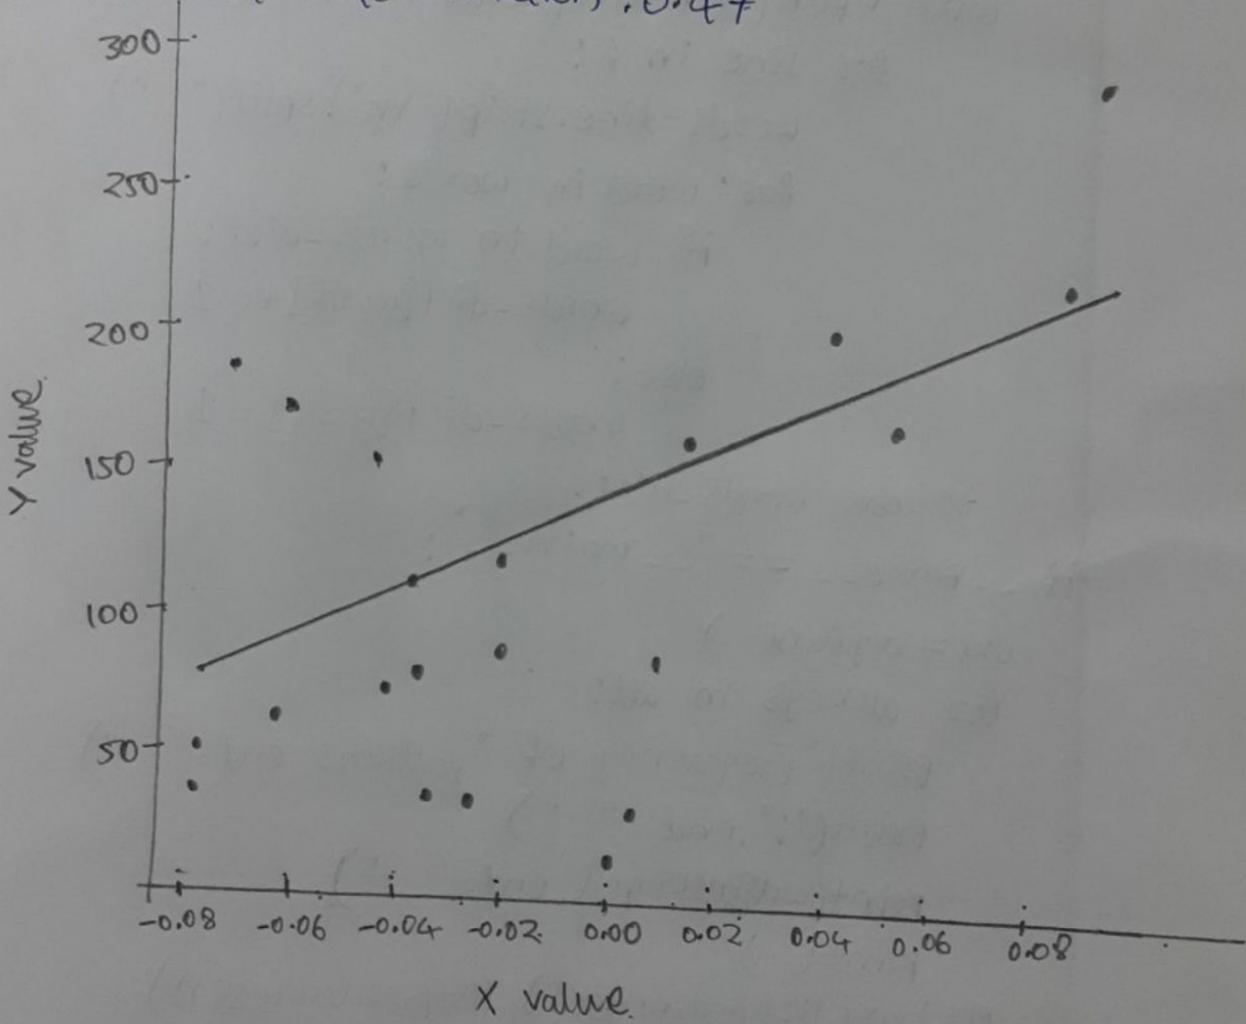
Constant : 152.918861826

Mean squared log error: 0.22

Mean squared error: 2548.07

Mean absolute error: 41.23

Coefficient of determination : 0.47



TITLE: Linear Regression -1

DATE: 11-4-22

Aim: Implement a program for Linear Regression with a sklearn dataset

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
from sklearn import datasets, linear_model
```

```
from sklearn.metrics import accuracy_score
```

```
from sklearn.metrics import mean_squared_error,  
r2_score, mean_absolute_error, mean_squared_log_error
```

```
diabetes_x, diabetes_y = datasets.load_diabetes(return_X_y  
=True)
```

```
data = datasets.load_diabetes(as_frame=True)
```

```
df = pd.DataFrame(data=data.data, columns=data.  
feature_names)
```

```
print(data.feature_names)
```

```
diabetes_x = diabetes_x[:, np.newaxis, 2]
```

```
diabetes_x_train = diabetes_x[:-20]
```

```
diabetes_x_test = diabetes_x[-20:]
```

```
diabetes_y_train = diabetes_y[:-20]
```

```
diabetes_y_test = diabetes_y[-20:]
```

```
regressor = linear_model.LinearRegression()
```

```
regressor.fit(diabetes_x_train, diabetes_y_train)
```

```
diabetes_y_pred = regressor.predict(diabetes_x_test)
```

```
print("Coefficient : ", regressor.coef_)
```

```
print("Const:", reg.s.intercept_)

print("Mean squared log error: %.2f" % mean_squared_
      log_error(diabetes_y-test, diabetes_y-pred))

print("Mean squared error: %.2f" % mean_squared_
      error(diabetes_y-test, diabetes_y-pred))

print("Mean Absolute error: %.2f" % mean_absolute_
      error(diabetes_y-test, diabetes_y-pred))

print("Coefficient of determination: %.2f"
      % r2_score(diabetes_y-test, diabetes_y-pred))

plt.scatter(diabetes_x-test, diabetes_y-test,
            color="black")

plt.plot(diabetes_x-test, diabetes_y-pred, color="blue",
         linewidth=2)

plt.xlabel("X value")
plt.ylabel("Y Value")
plt.xticks()
plt.yticks()

plt.savefig("W2a-plot.png", format="png")
plt.show()
```

Output -

[0.75362795])

Linear Regression ()

[9054.91] [20885.] [18265.6293]

[9437.37] [16919.] [18553.8618]

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

,

TITLE: Linear Regression - 2

DATE: 18-4-22

Aim: Implement a program for Linear Regression by reading CSV files

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn.linear_model

def prepare_country_stats(oecd_bli, gdp_per_capita):
    oecd_bli = oecd_bli[oecd_bli["INEQUALITY"] != "TOT"]
    oecd_bli = oecd_bli.pivot(index="country",
                               columns="Indicators", values="value")
    gdp_per_capita = gdp_per_capita.rename(columns={"2015": "GDP per capita"}, inplace=True)
    gdp_per_capita.set_index("country", inplace=True)
    full_country_stats = pd.merge(left=oecd_bli,
                                  right=gdp_per_capita, left_index=True,
                                  right_index=True)
    full_country_stats.sort_values(by="GDP per capita", inplace=True)

    remove_indices = [0, 1, 6, 8, 33, 34, 35]
    keep_indices = list(set(range(36)) - set(remove_indices))
    return full_country_stats[["GDP per capita", "Housing expenditure"]].iloc[keep_indices]

oecd_bli = pd.read_csv("../datasets/oecd-bli-2015.csv",
                      thousands=',')

```

```

gdp_per_capita = pd.read_csv("../datasets/gdp_per_capita.csv",
                             thousands=',')
country_stats = prepare_country_stats(oecd_bli,
                                       gdp_per_capita)
country_stats.to_csv("gdp.csv", sep=";")

X = np.c_[country_stats["GDP per capita"]]
y = np.c_[country_stats["Housing expenditure"]]

model = sklearn.linear_model.LinearRegression()
model.fit(X, y)

X_new = [[32587]]
print(model.coef_)
print(model)

y_pred = model.predict(X)
for i in range(X.size):
    print(X[i], y[i], y_pred[i])

plt.scatter(X, y)
plt.xlabel("GDP per capita")
plt.ylabel("Housing expenditure")
plt.plot(X, y_pred)
plt.savefig("../datasets/gdp_per_capita_personal_earnings.png")
plt.show()

```

Output -

Coefficients:

{[0.000732]}

{-0.064531}]

1

$$[-0.047356]$$

constant;

$[-0.550179 \quad -0.550153 \quad -\quad -\quad -]$

-0.550159]

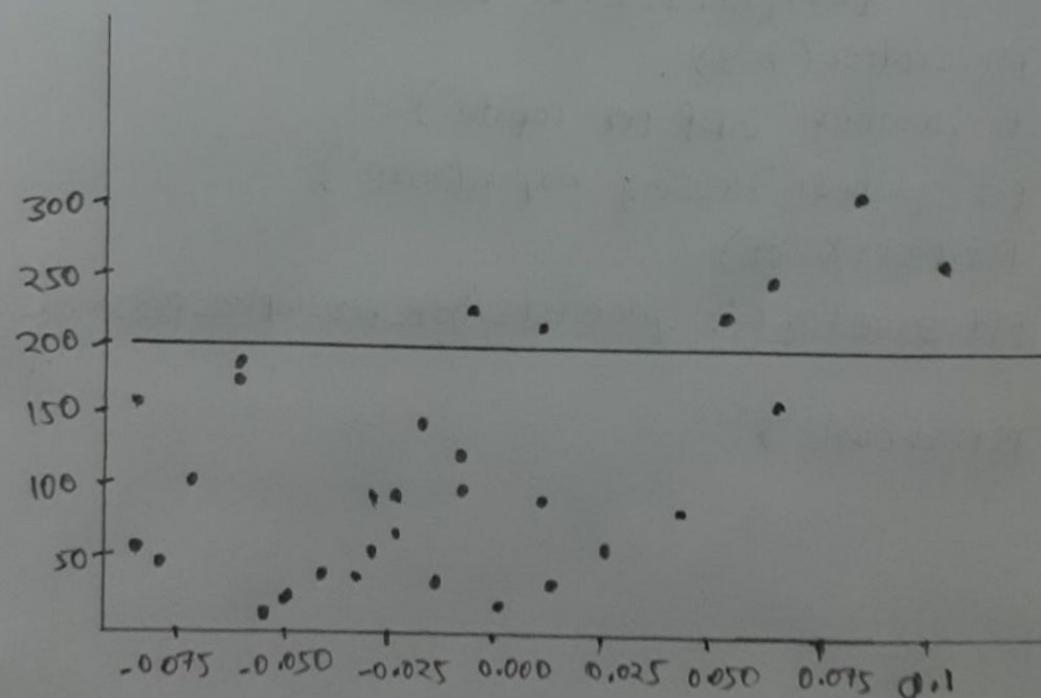
Mean Squared error: 9601.1

Coefficient of determination: -0.86.

{261. 113. 131. — — .

— — — — 200. 200.]

{[0 0 0 0 0 0 0 1 0 0 0] - ____.]



TITLE: Logistic Regression - 1

DATE: 25-4-22

Aim: Implement a program for Logistic Regression with any sklearn dataset

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn import datasets, linear_model
```

```
from sklearn.metrics import mean_squared_error,  
r2_score, confusion_matrix.
```

```
diabetes_X, diabetes_y = datasets.load_diabetes(return_X_y  
=True)
```

```
print(type(diabetes_X))
```

```
diabetes_X = diabetes_X[:, np.newaxis]
```

```
diabetes_X_train = diabetes_X[:-30]
```

```
diabetes_X_test = diabetes_X[-30:]
```

```
diabetes_y_train = diabetes_y[:-30]
```

```
diabetes_y_test = diabetes_y[-30:]
```

```
regressor = linear_model.LogisticRegression()
```

```
regressor.fit(diabetes_X_train, diabetes_y_train)
```

```
diabetes_y_pred = regressor.predict(diabetes_X_test)
```

```
print("Coefficients: \n", regressor.coef_)
```

```
print("Constant: \n", regressor.intercept_)
```

```
print("Mean Squared Error: %.2f" % mean_squared_error(  
diabetes_y_test, diabetes_y_pred))
```

```
print("Coefficient of Determination: %.2f" % r2_score(diabetes_y_test, diabetes_y_pred))  
plt.scatter(diabetes_X_test, diabetes_y_test, color="black")  
plt.plot(diabetes_X_test, diabetes_y_pred, color="blue")  
plt.xticks()  
plt.yticks()  
print(diabetes_y_test, diabetes_y_pred)  
print(confusion_matrix(diabetes_y_test, diabetes_y_pred))  
plt.show()
```

Output -

$\{ -3.104213 \times 10^{-5} \}$

$\{ \}$

$\{ 1.6358112 \times 10^{-5} \}$

LogisticRegression()

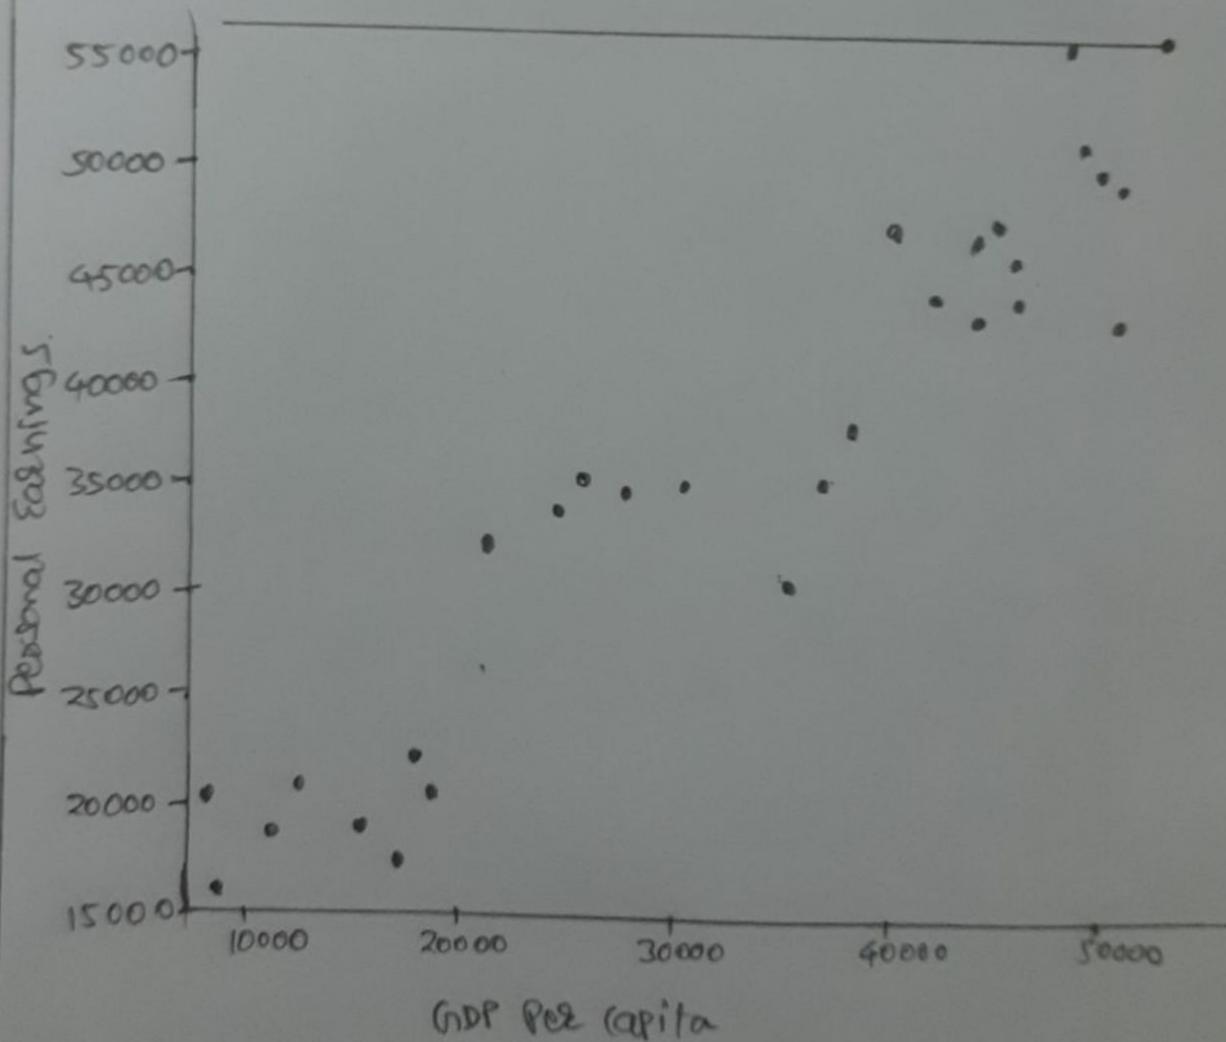
$\{ 9054.91 \} \quad \{ 20885. \} \quad \{ 56340.0 \}$

$\{ 9437.37 \} \quad \{ 16919. \} \quad 56340.0$

$\{ \}$

$\{ \}$

$\{ 55805.2 \} \quad \{ 56340. \} \quad 56340.0$



TITLE: Logistic Regression - 2

DATE: 30-5-22

Aim: Implement a program for Logistic Regression by reading csv files

```

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import sklearn.linear_model
def prepare_country_stats(oecd_bli, gdp_per_capita):
    oecd_bli = oecd_bli[oecd_bli["INEQUALITY"] == "TOT"]
    oecd_bli = oecd_bli.pivot(index="Country",
                               columns="Indicator", values="Value")
    gdp_per_capita.rename(columns={"2015": "GDP Per Capita"}, inplace=True)
    gdp_per_capita.set_index("Country", inplace=True)
    full_country_stats = pd.merge(left=oecd_bli,
                                  right=gdp_per_capita,
                                  left_index=True, right_index=True)
    remove_indices = [0, 1, 6, 8, 33, 34, 35]
    keep_indices = list(set(range(36)) - set(remove_indices))
    return full_country_stats[["GDP Per Capita",
                               "Personal earnings"]].iloc[keep_indices]
oecd_bli = pd.read_csv("../datasets/oecd-bli-2015.csv",
                      thousands=',')

```

```

gdp_per_capita = pd.read_csv("../datasets/gdp_per_capita.csv", thousands=',')
country_stats = prepare_country_stats(oecd_bli, gdp_per_capita)

country_stats.to_csv("gdp.csv", sep=";")

X = np.c_[country_stats[["GDP per Capita"]]]
y = np.c_[country_stats[["Personal Earnings"]]]

print(y.size)

y = y.reshape(y.size,)

model = sklearn.linear_model.LogisticRegression()

model.fit(X, y)

X_new = [[22587]]
print(model.coef_)

print(model.estimator_type)

y_pred = model.predict(X)

for i in range(X.size):
    print(X[i], y[i], y_pred[i])

plt.scatter(X, y)

plt.xlabel("GDP Per Capita")
plt.ylabel("Personal Earnings")
plt.plot(X, y_pred)

plt.savefig("../datasets/gdp_per_capita_Personal_Earnings.png", dpi=600)

plt.show()

```

Output -

$$y = 4.942 + -2.227 \times x_0 + 2.250 \times x_1$$

$$y = 6.179 + -2.069 \times x_0 + 1.723 \times x_1$$

$$y = 12.355 + -1.601 \times x_0 + -0.796 \times x_1$$

Support Vectors

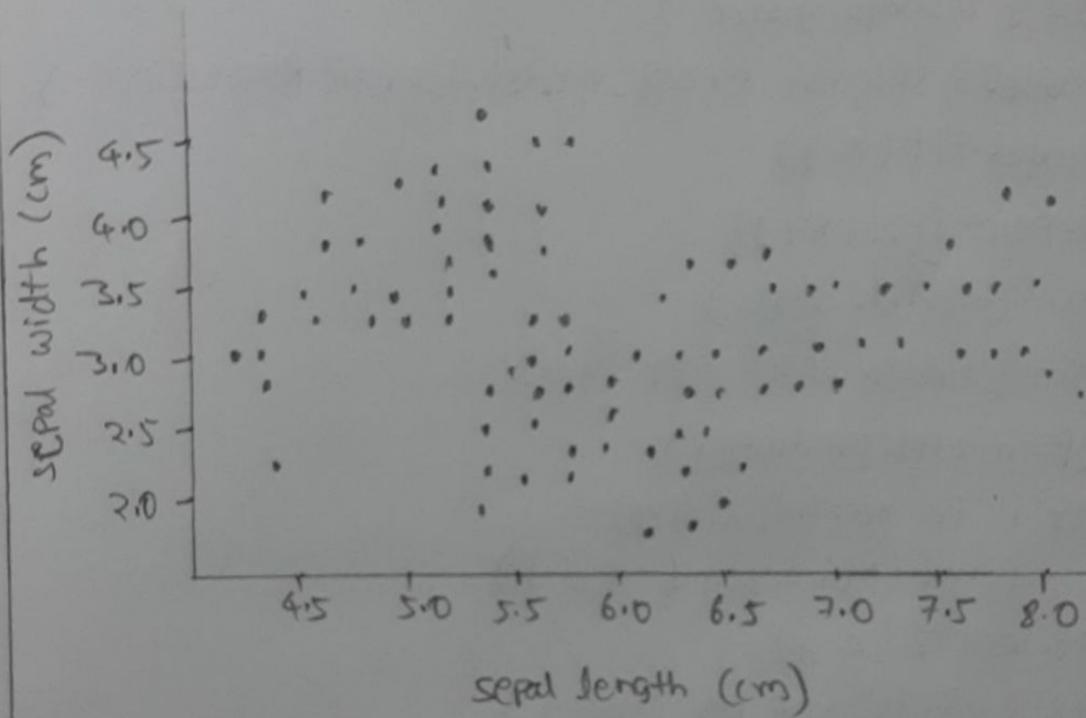
$$\begin{bmatrix} 4.9 & 3.0 \end{bmatrix}$$

$$\begin{bmatrix} 5.7 & 3.8 \end{bmatrix}$$

1

1

$$\begin{bmatrix} 5.9 & 3.7 \end{bmatrix}$$



TITLE: SVM - 1

DATE: 6-6-22

Aim: Write a program to construct a SVM classifier with a sklearn dataset

```
import numpy as np
```

```
from sklearn import datasets
```

```
from sklearn.pipeline import Pipeline
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.svm import LinearSVC
```

```
from sklearn import svm
```

```
import matplotlib.pyplot as plt
```

```
import pandas as pd
```

```
import sklearn.metrics as skm
```

```
iris = datasets.load_iris()
```

```
x = iris["data"][:, (0, 1)]
```

```
y = (iris["target"]).astype(np.float64)
```

```
svm_clf = svm.SVC(kernel="linear")
```

```
svm_clf.fit(x, y)
```

```
for (intercept, coef) in zip(svm_clf.intercept_, svm_clf.coef_):
```

```
S = "y = {} + {}x_1".format(intercept, coef)
```

```
for (i, c) in enumerate(coef):
```

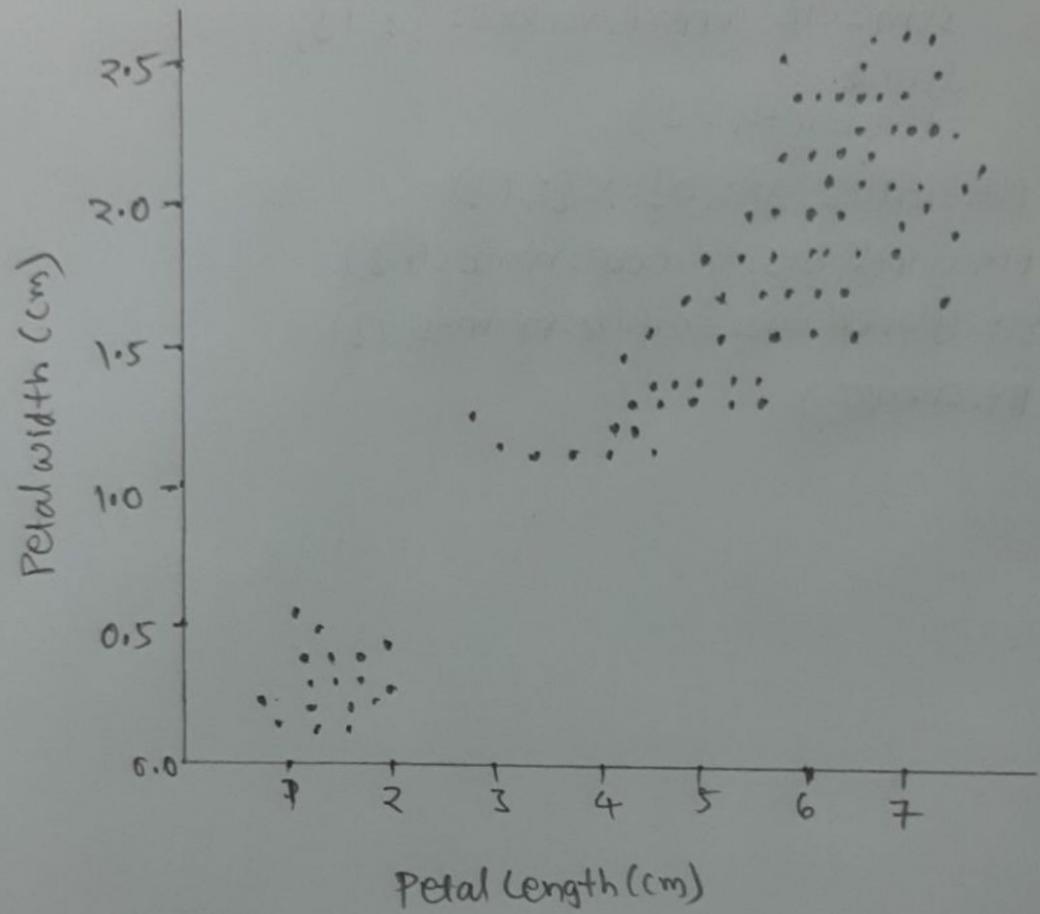
```
S += "+ {}x_2".format(c, i)
```

```
print(S)
```

```
print(x)
```

```
point("Support Vectors")
point(svm-clf, support-vectors-)
plt.scatter(
    svm-clf.support_vectors_[:,0],
    svm-clf.support_vectors_[:,1],
    linear
    linewidth=1)
plt.scatter(x[:,0], x[:,1])
plt.xlabel(iris.feature_names[0])
plt.ylabel(iris.feature_names[1])
plt.show()
```

Output -
array ([1])



TITLE: SVM - 2

DATE: 13-6-22

Aim: Write a program to construct a SVM classifier

```

import numpy as np
from sklearn import datasets
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.svm import LinearSVC
import matplotlib.pyplot as plt
import pandas as pd

iris = datasets.load_iris()
irisdf = pd.DataFrame(data=np.c_[iris['data'], iris['target']],
columns=iris['feature_names'] + ['target'])

irisdf.plot.scatter(x="Petal length (cm)", y="Petal width (cm)")

svm_clf = Pipeline([
    ("scaler", StandardScaler()),
    ("linear_svc", LinearSVC(C=1, loss="hinge")),
])

```

svm_clf.fit(x, y)

print(svm_clf.predict([[5.5, 1.7]]))

$x_1, y_1 = [1, 7], [1, 0.75]$

$x_2, y_2 = [2.5, 2], [3, 0]$

$x_3, y_3 = [1.5, 3], [2.5, 0]$

```
plt.plot(x1, y1, x2, y2, x3, y3, marker = 'o')  
plt.scatter(x[:, 0], x[:, 1])  
plt.xlabel(iris.feature_names[2])  
plt.ylabel(iris.feature_names[3])  
plt.show()
```

Output -

array([1])

Training Accuracy: 0.969230

Testing Accuracy: 0.9

TITLE: Decision Tree

DATE: 20-6-22

Aim: Write a program to demonstrate working of Decision Tree

```

from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier
import matplotlib.pyplot as plt
iris = load_iris()
X_train = iris.data[:-20, 2:] # petal length and width
y_train = iris.target[:-20]
X_test = iris.data[-20:, 2:]
y_test = iris.target[-20:]
tree_clf = DecisionTreeClassifier(max_depth=2, random_state=42)
tree_clf.fit(X_train, y_train)
tree_clf.predict([[5, 1.5]])
print("Training Accuracy:", tree_clf.score(X_train, y_train))
print("Testing Accuracy:", tree_clf.score(X_test, y_test))
y_pred = tree_clf.predict(X_test)
for i in range(0, 20):
    print(y_test[i], " ", y_pred[i])

```

